Recoverability

In database management, the terms "**recoverability**," "recoverable schedule," and "**irrecoverable schedule**" pertain to the ability of a database system to ensure that transactions can be recovered and maintain data consistency, even in the presence of failures.

1. Recoverability:

Definition:
It is a property of a schedule in a database system that ensures that, in the event of a failure (such as a system crash), the system can recover to a consistent state.

Objective:
The primary goal is to prevent situations where a transaction reads data that has been modified by another transaction that later gets rolled back due to a failure.

Key Requirement:
A schedule is recoverable if, for each pair of transactions T1 and T2, where T1 writes data that T2 reads, T1 must commit before T2 starts.

2. Recoverable Schedule:

Definition:
A schedule is considered recoverable if it adheres to the recoverability property.
In other words, it ensures that no transaction reads uncommitted data from another transaction.

Guarantee:
In a recoverable schedule, if a transaction T1 reads data written by another transaction T2, it is guaranteed that T2 has already committed when T1 starts.

Ensuring Recoverability:
To achieve recoverability, a database system can use techniques like locking, timestamps, or validation checks to ensure that transactions do not read uncommitted data.

3. Irrecoverable Schedule:

Definition:
It is a schedule that does not adhere to the recoverability property.
It allows transactions to read data that has not been committed by other transactions.

Problem:
They can lead to data inconsistencies, as a transaction might read changes that are later rolled back, causing incorrect results or violating the isolation property.

Risk:
They are typically considered undesirable because they can lead to unpredictable and incorrect behavior in the presence of failures.

Summary:
- Recoverability is a crucial concept in database systems to ensure data consistency and integrity, especially in the event of system failures.
- A recoverable schedule guarantees that transactions do not read uncommitted data from other transactions and, therefore, helps maintain the ACID (Atomicity, Consistency, Isolation, Durability) properties of the database.
- Conversely, an irrecoverable schedule poses a risk of data inconsistencies and is generally avoided in database design and transaction scheduling.
- Techniques like locking and timestamp-based concurrency control are commonly used to enforce recoverability in database systems.

In transaction management, particularly in the context of database systems, the terms "cascading," "cascadeless," and "strict cascadeless" schedules are used to describe different approaches to handling the effects of transactions.

These concepts relate to how changes made by one transaction affect other transactions in the system.

Cascading Schedule:

- In a cascading schedule, the effects of one transaction can affect other transactions in the system.

- Specifically, if one transaction makes changes to the database and then another transaction relies on those changes, the second transaction may also be affected by the changes made by the first transaction.

- This can lead to a domino effect where the failure or rollback of one transaction can impact subsequent transactions that depended on its changes.

Cascadeless Schedule:

- It is designed to prevent the cascading effects of one transaction's changes on others.

- In a cascadeless schedule, a transaction is only allowed to read data that was committed by other transactions.

- It cannot read uncommitted data.

- It ensures that a transaction's outcome is not influenced by other transactions that are still in progress or that later get rolled back.

Strict Cascadeless Schedule:

- It is a more stringent version of a cascadeless schedule.

- In addition to not allowing transactions to read uncommitted data, It also prevents transactions from writing data that might be read by other transactions before they are committed.

- This ensures that no transaction can have its outcome affected by changes made by other transactions, even if they commit their changes before the transaction in question.

- It provide the highest level of isolation among transactions.

- They are often preferred in systems where data consistency and integrity are critical.

- They can lead to increased contention and potentially reduce concurrency, as transactions may need to wait for exclusive access to data.

- The choice of which schedule to use depends on the specific requirements of the application and the trade-offs between data consistency, concurrency, and isolation.

- Different database systems may offer different isolation levels to accommodate these needs, such as READ COMMITTED, REPEATABLE READ, and SERIALIZABLE, each with varying degrees of strictness in their scheduling rules.