

Concurrency control problem in more detail with examples:

### **Lost Updates:**

Problem:

When two or more transactions try to update the same data simultaneously, one transaction's updates may be lost if they are overwritten by another transaction's updates.

**Example:** Consider two bank transactions:

Transaction 1: Transfer \$100 from Account A to Account B.

Transaction 2: Transfer \$150 from Account A to Account B.

- If both transactions run concurrently and read the balance of Account A (e.g., \$500) at the same time, they may both calculate the new balance incorrectly.
- Without proper concurrency control, one of these transactions might update the balance to \$350, and the other may update it to \$400, causing a lost update.

### **Dirty Reads:**

Problem:

A dirty read occurs when one transaction reads data that has been modified by another transaction but not yet committed.

**Example:** Transaction 1 updates a customer's address:

- Transaction 1: Update Customer A's address to "New Address."
- Transaction 2 reads Customer A's data before Transaction 1 commits.
- If Transaction 1 rolls back for some reason after Transaction 2 has read the "New Address," Transaction 2 will have read incorrect and uncommitted data.

### **Non-Repeatable Reads:**

Problem:

Inconsistent data can be read during multiple reads by a transaction if another transaction modifies the data in between.

Example: Transaction 1 reads the price of a product:

- Transaction 1: Read Product X's price (e.g., \$100).
- Transaction 2 updates Product X's price to \$120.

- Transaction 1 reads Product X's price again.
- Transaction 1 will observe a different price during the second read, resulting in non-repeatable reads.

## **Phantom Reads:**

Problem:

A transaction reads a set of rows that satisfy a condition, but before the transaction can complete, another transaction inserts, updates, or deletes rows that match the condition.

Example: Transaction 1 queries all products with a price less than \$50:

- Transaction 1: `SELECT FROM Products WHERE Price < $50;`
- Transaction 2 inserts a new product with a price of \$40.
- Transaction 1 re-executes the same query.
- Transaction 1 will see a different set of rows in the second query, including the newly inserted product, causing phantom reads.

## **Deadlocks:**

Problem:

Deadlocks occur when two or more transactions are waiting for each other to release locks on resources, leading to a system stall.

Example: Consider two transactions:

- Transaction 1: Locks Resource A, requests Resource B.
- Transaction 2: Locks Resource B, requests Resource A.
- If both transactions acquire their initial locks and then request the opposite resource, they will enter a deadlock state, as neither can proceed without the other releasing its lock.
- To mitigate these problems, DBMSs use various concurrency control mechanisms such as locking, timestamps, and isolation levels (e.g., Read Committed, Serializable).
- These mechanisms ensure that transactions maintain data consistency and integrity while allowing for concurrent execution.
- Properly configuring these mechanisms and testing your application's behavior under concurrent conditions is crucial for avoiding these concurrency control problems.

